

Car Price Prediction

Linear Regression

Index

- Reading and Understanding data
- Data cleaning and Preparation
- Visualizing the data
- Deriving new features
- Bivariate analysis
- Creating dummy variable
- Splitting the train-test and Rescaling the features
- Building the model
- Residual analysis on train set
- Prediction and evaluation

Step 1: Reading and Understanding data

```
In [83]: df.columns
```

```
Out[83]: Index(['car_ID', 'symboling', 'CarName', 'fueltype', 'aspiration',  
              'doornumber', 'carbody', 'drivewheel', 'enginelocation', 'wheelbase',  
              'carlength', 'carwidth', 'carheight', 'curbweight', 'enginetype',  
              'cylindernumber', 'enginesize', 'fuelsystem', 'boreratio', 'stroke',  
              'compressionratio', 'horsepower', 'peakrpm', 'citympg', 'highwaympg',  
              'price'],  
             dtype='object')
```

```
| In [84]: df.shape
```

```
Out[84]: (205, 26)
```

Reading the csv file and creating data Frame (df)

Step 2: Data cleaning and Preparation

```
In [87]: # splitting company name from CarName
# dropping CarName column

df['Company'] = df['CarName'].apply(lambda x: x.split()[0])
df.drop(['CarName'], axis=1, inplace=True)
df.head()
```

Out[87]:

	car_ID	symboling	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocation	wheelbase	carlength	...	fuelsystem	boreratio	stroke	compre
0	1	3	gas	std	two	convertible	rwd	front	88.6	168.8	...	mpfi	3.47	2.68	
1	2	3	gas	std	two	convertible	rwd	front	88.6	168.8	...	mpfi	3.47	2.68	
2	3	1	gas	std	two	hatchback	rwd	front	94.5	171.2	...	mpfi	2.68	3.47	
3	4	2	gas	std	four	sedan	fwd	front	99.8	176.6	...	mpfi	3.19	3.40	
4	5	2	gas	std	four	sedan	4wd	front	99.4	176.6	...	mpfi	3.19	3.40	

5 rows x 26 columns

Splitting the Company from CarName variable

Spelling mistake correction

```
# correcting spelling mistakes of the Company
```

```
df.Company = df.Company.str.lower()
```

```
def replace_name(a,b):  
    df.Company.replace(a,b,inplace=True)
```

```
replace_name('maxda','mazda')
```

```
replace_name('porcshce','porsche')
```

```
replace_name('toyouta','toyota')
```

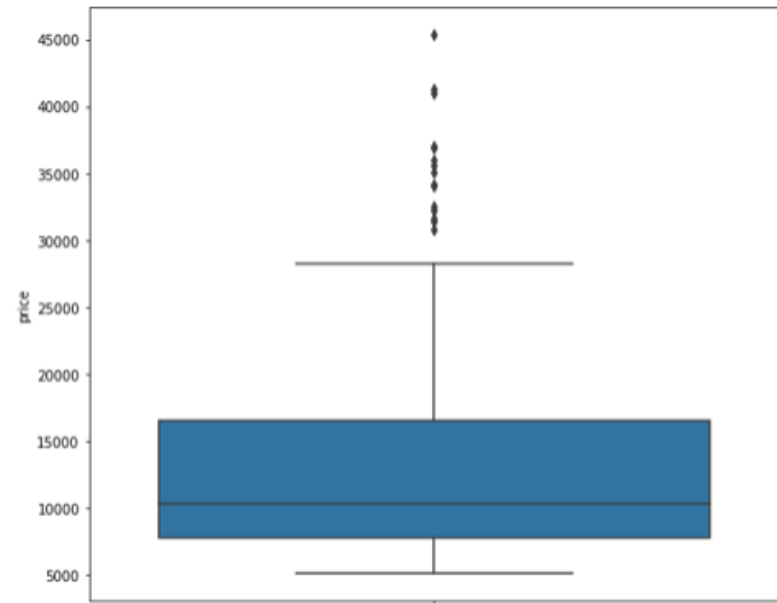
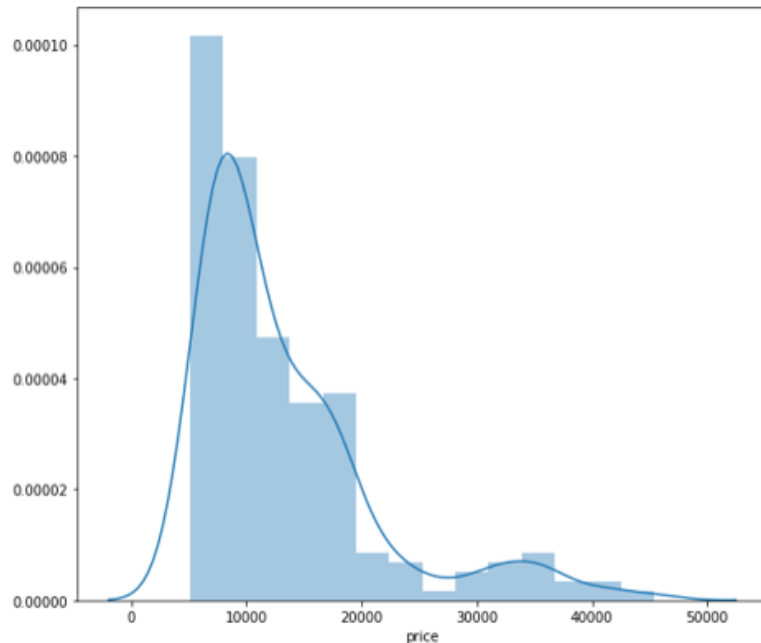
```
replace_name('vokswagen','volkswagen')
```

```
replace_name('vw','volkswagen')
```

```
df.Company.unique()
```

```
array(['alfa-romero', 'audi', 'bmw', 'chevrolet', 'dodge', 'honda',  
      'isuzu', 'jaguar', 'mazda', 'buick', 'mercury', 'mitsubishi',  
      'nissan', 'peugeot', 'plymouth', 'porsche', 'renault', 'saab',  
      'subaru', 'toyota', 'volkswagen', 'volvo'], dtype=object)
```

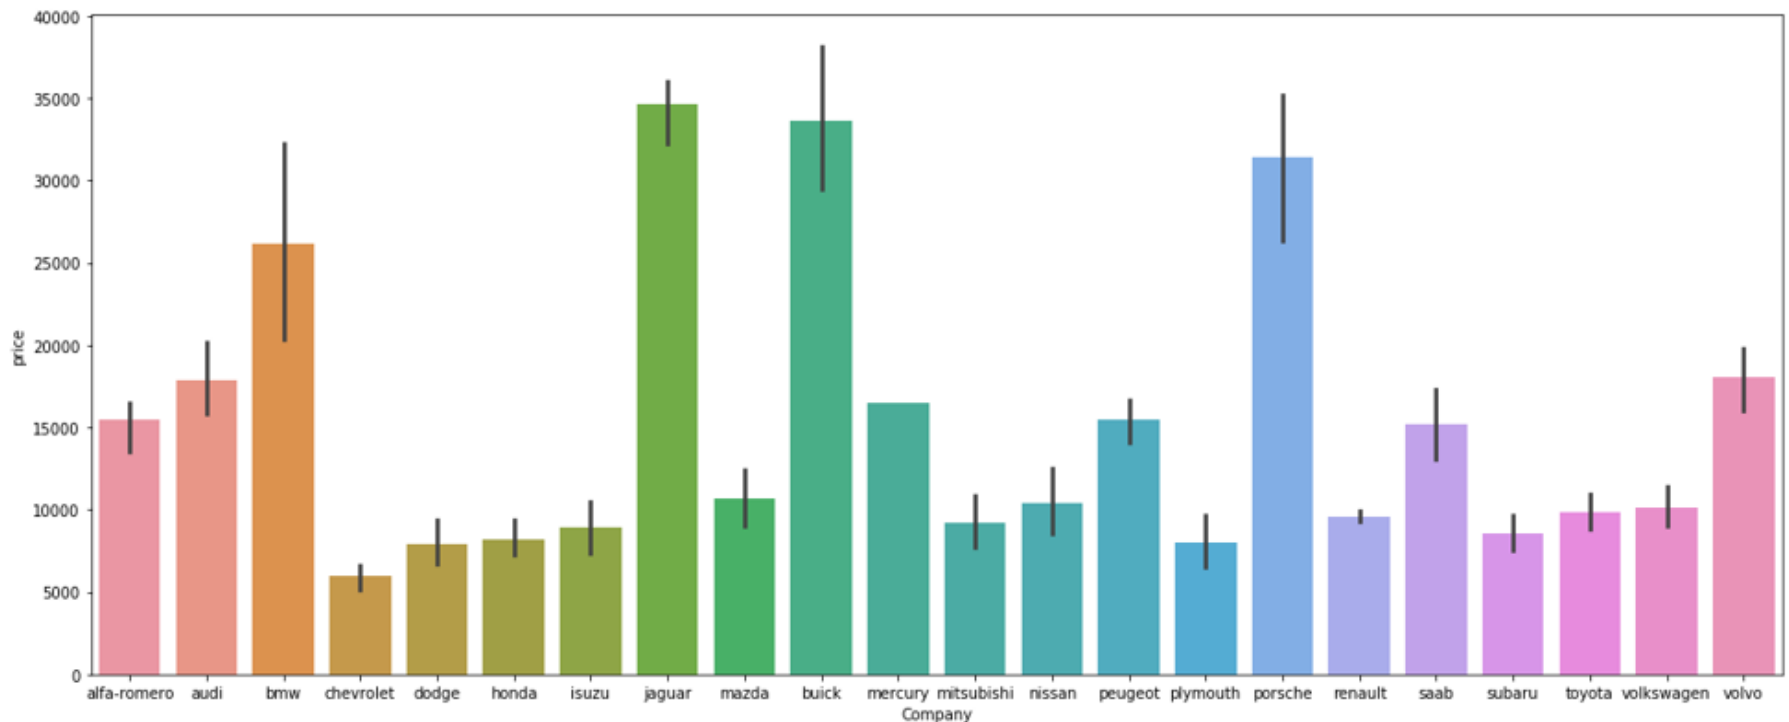
Step 3: Visualizing the data



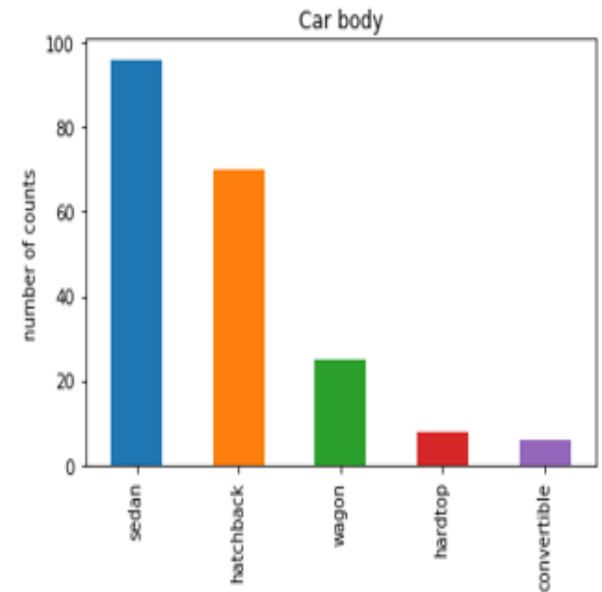
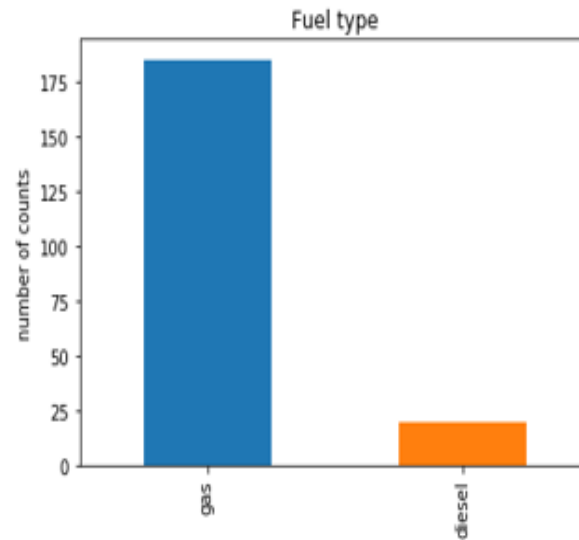
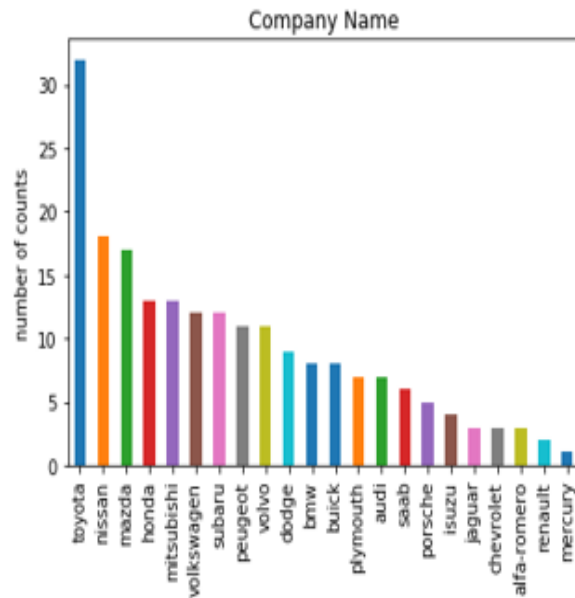
The graph is right skewed and there are not much outliers present in the price variable

- there is difference between mean and median of the price
 - mean price = 13276.710571
 - median price = 10295.000000
- the distribution of the price is rightly skewed
- the data points have high variance, 85% of the data has price below 18500 and 15% of the data has price between 22563 to 45400

Visualizing categorical data

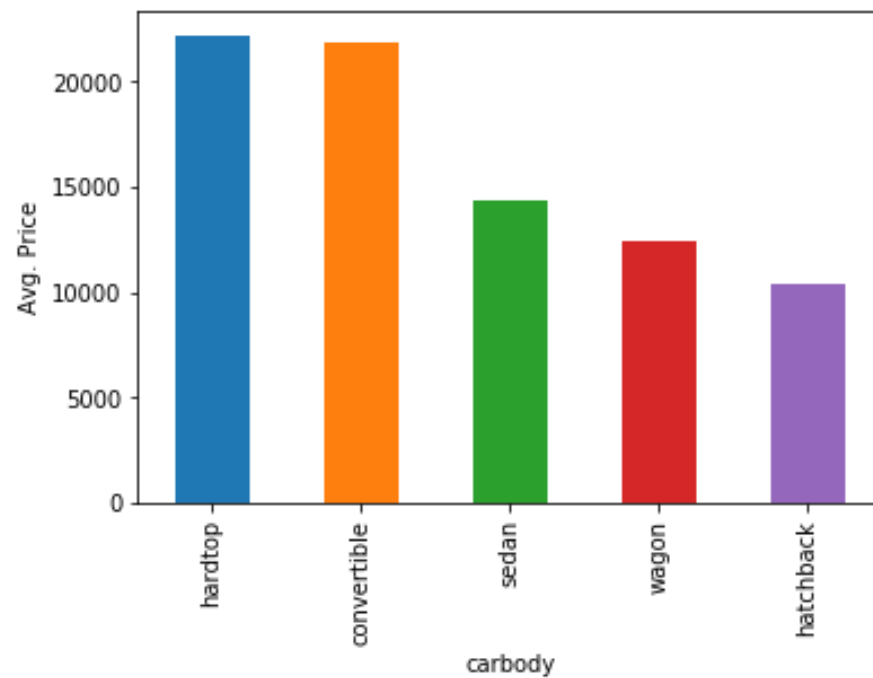
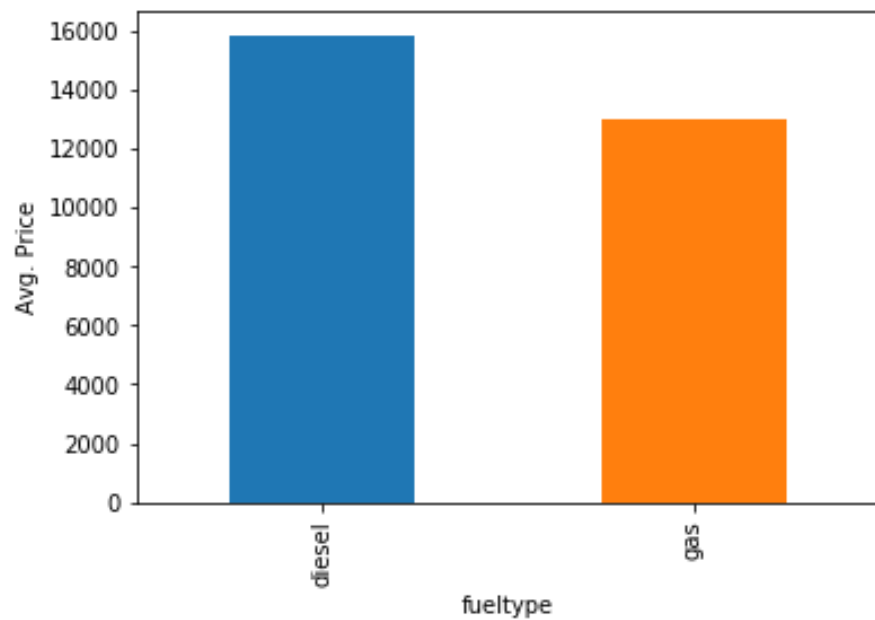


- Jaguar, Buick and porsche seem to be top 3 high price cars



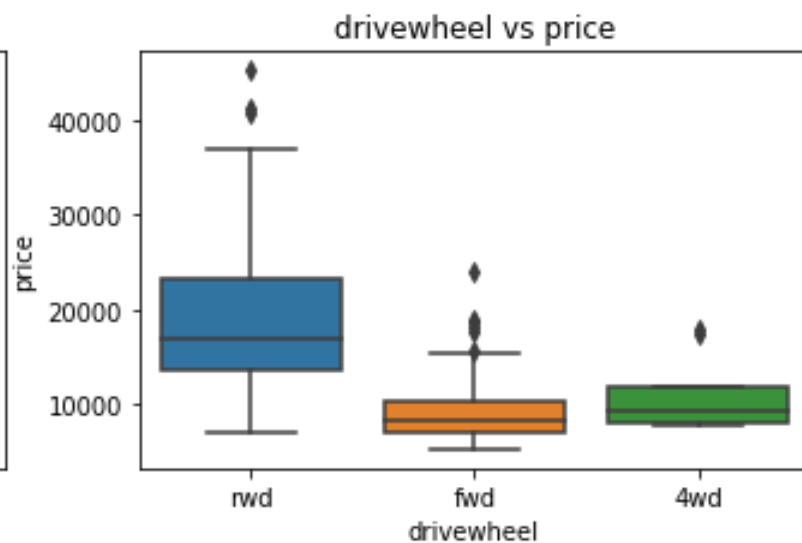
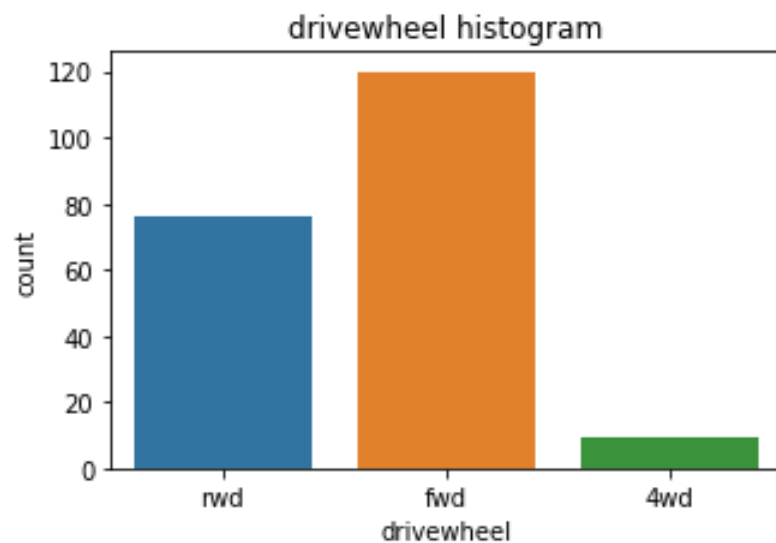
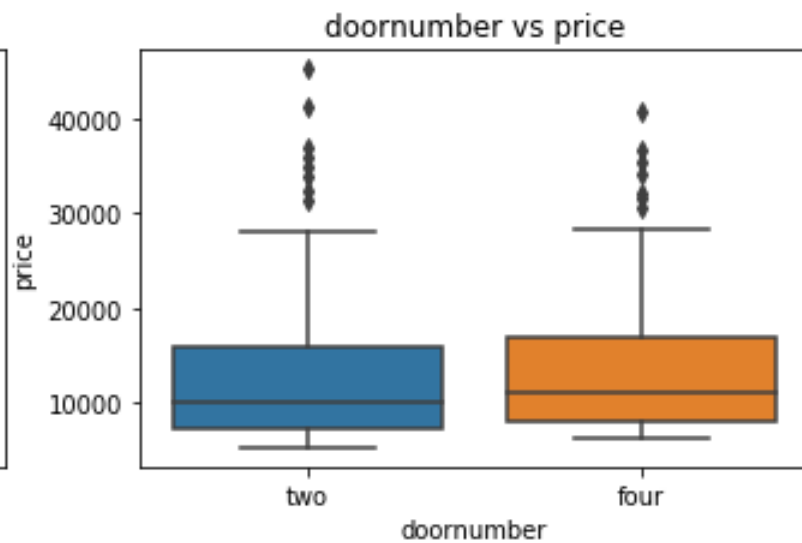
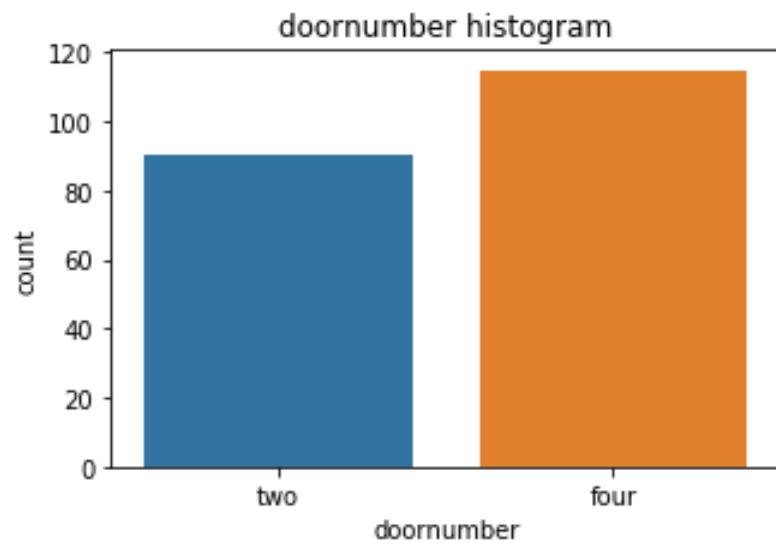
inference

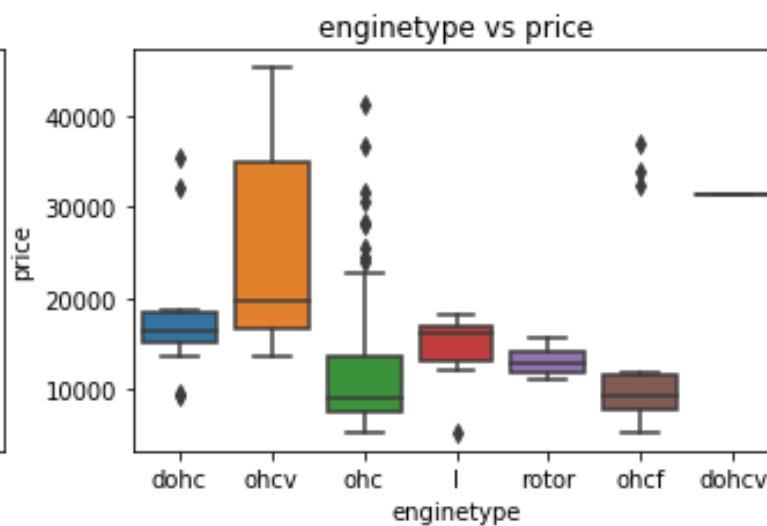
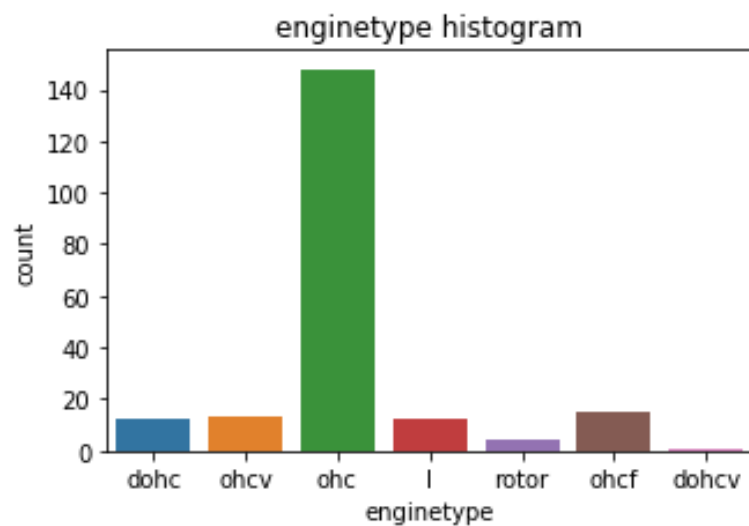
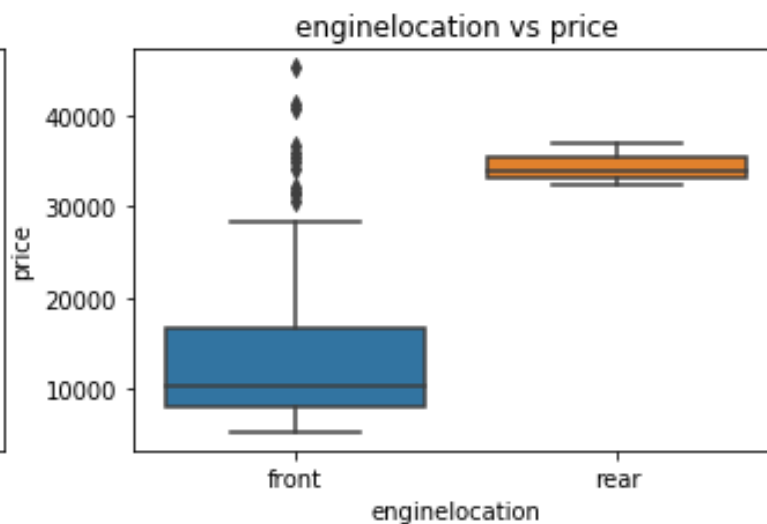
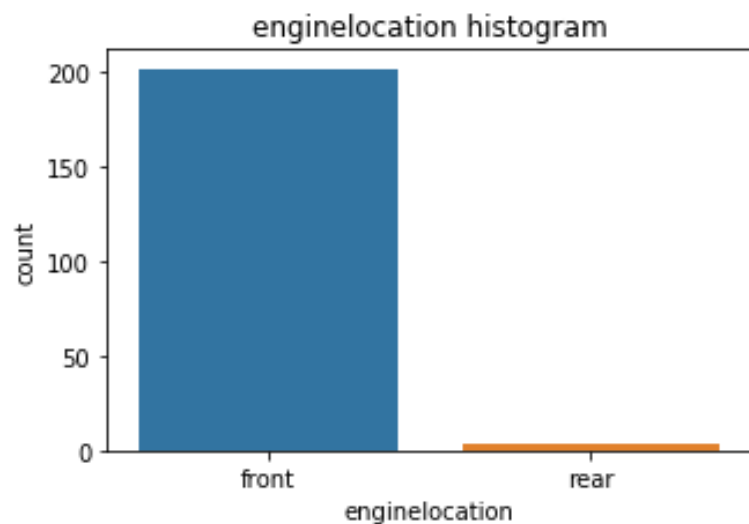
- toyota seem to be favoured company
- number of gas fuel type is greater than diesel fuel type car
- sedan is highly preferred



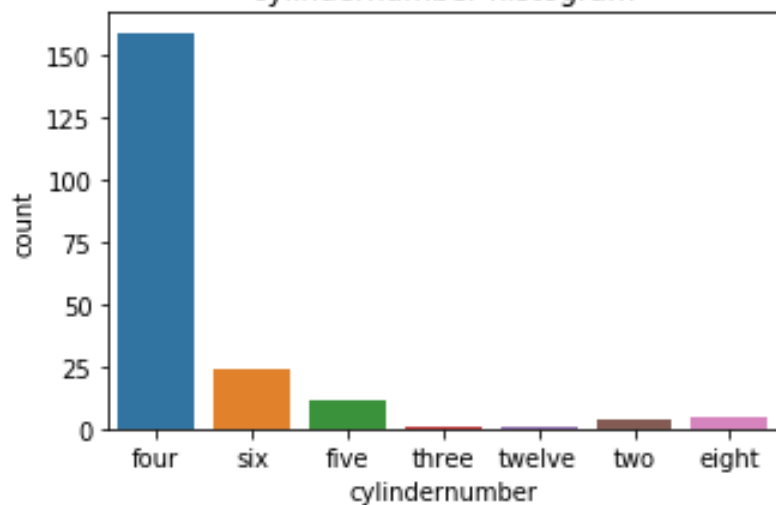
inference

- diesel car is more expensive than gas fuel type car
- hardtop and convertible cars are more expensive than the rest

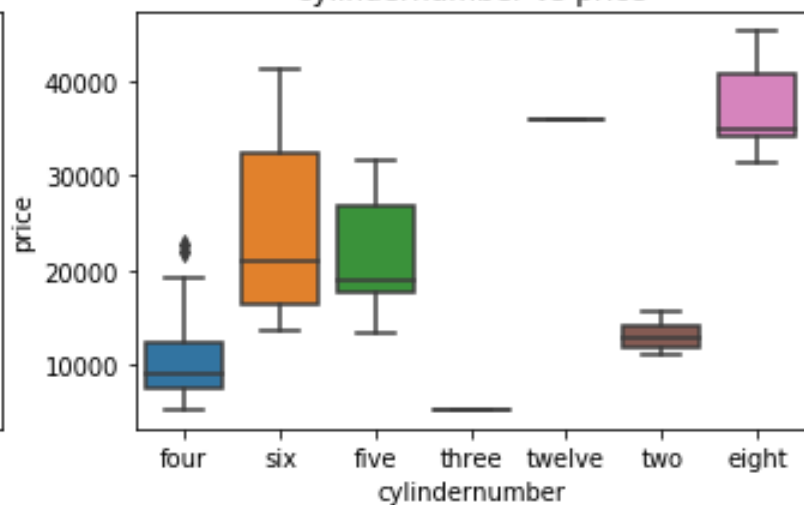




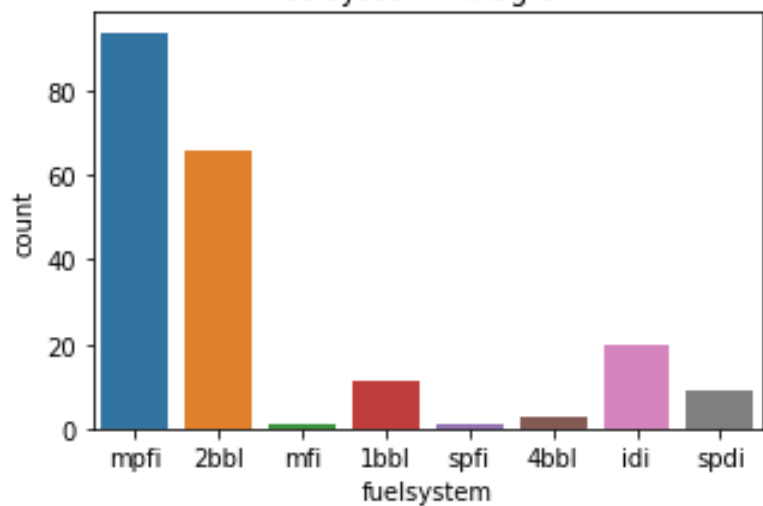
cylindernumber histogram



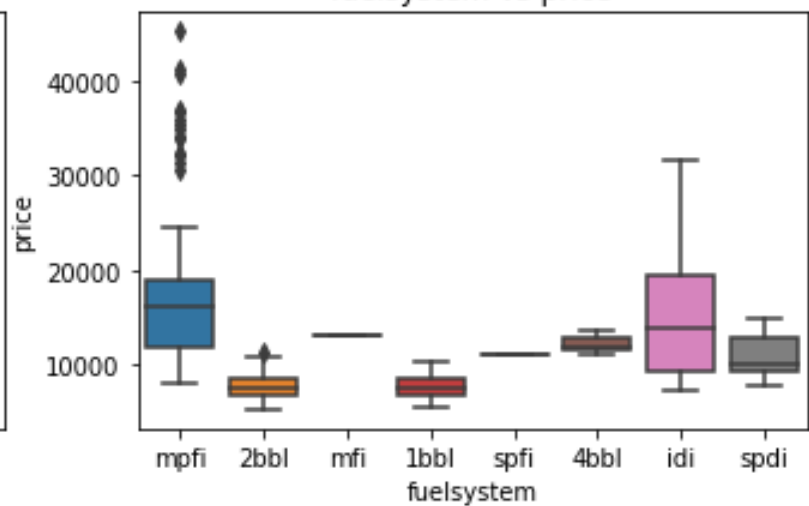
cylindernumber vs price

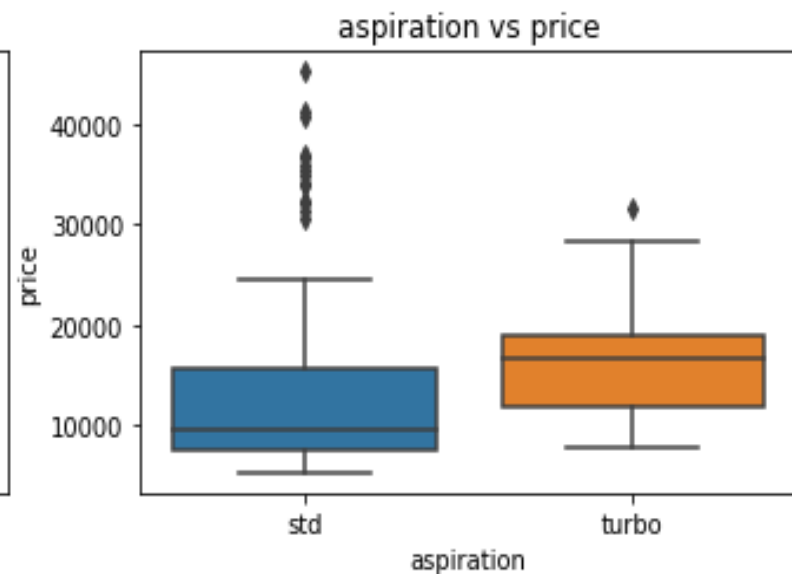
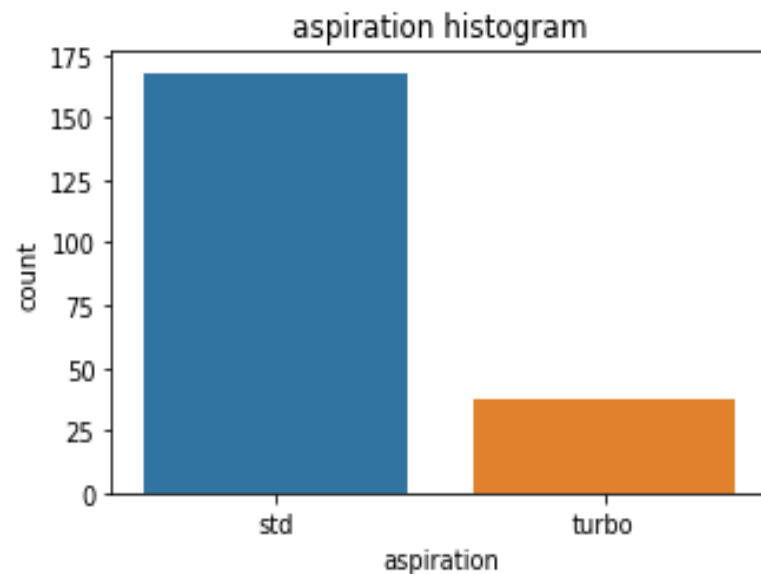


fuelsystem histogram



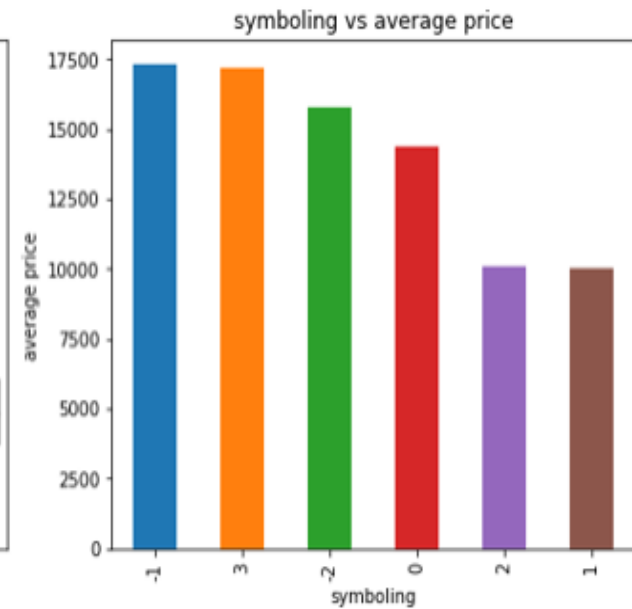
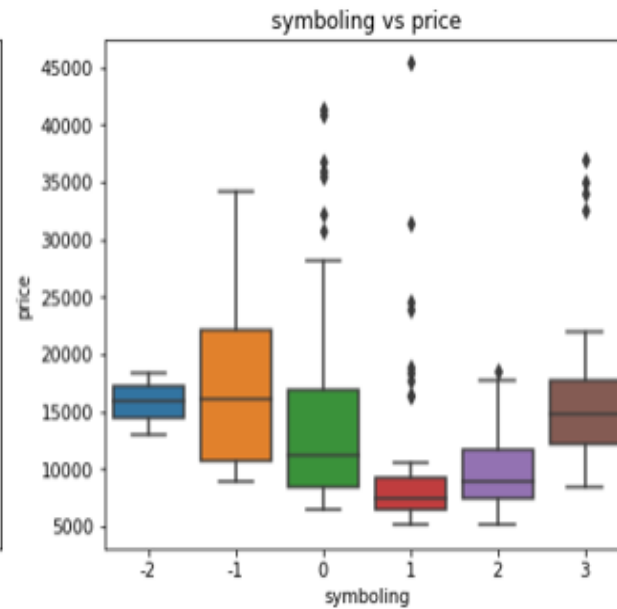
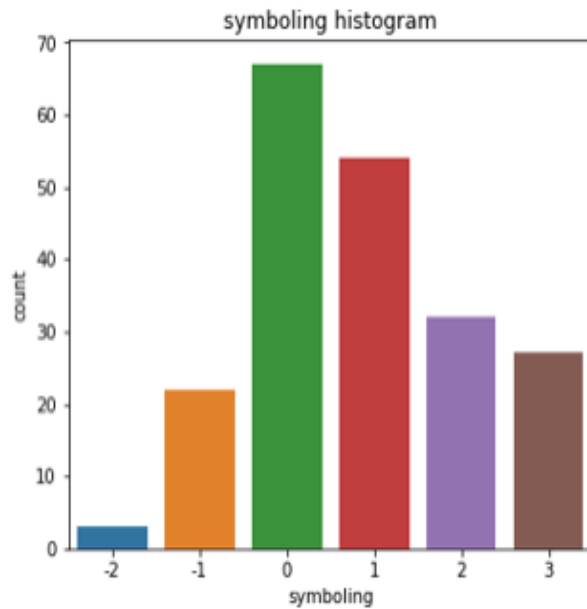
fuelsystem vs price





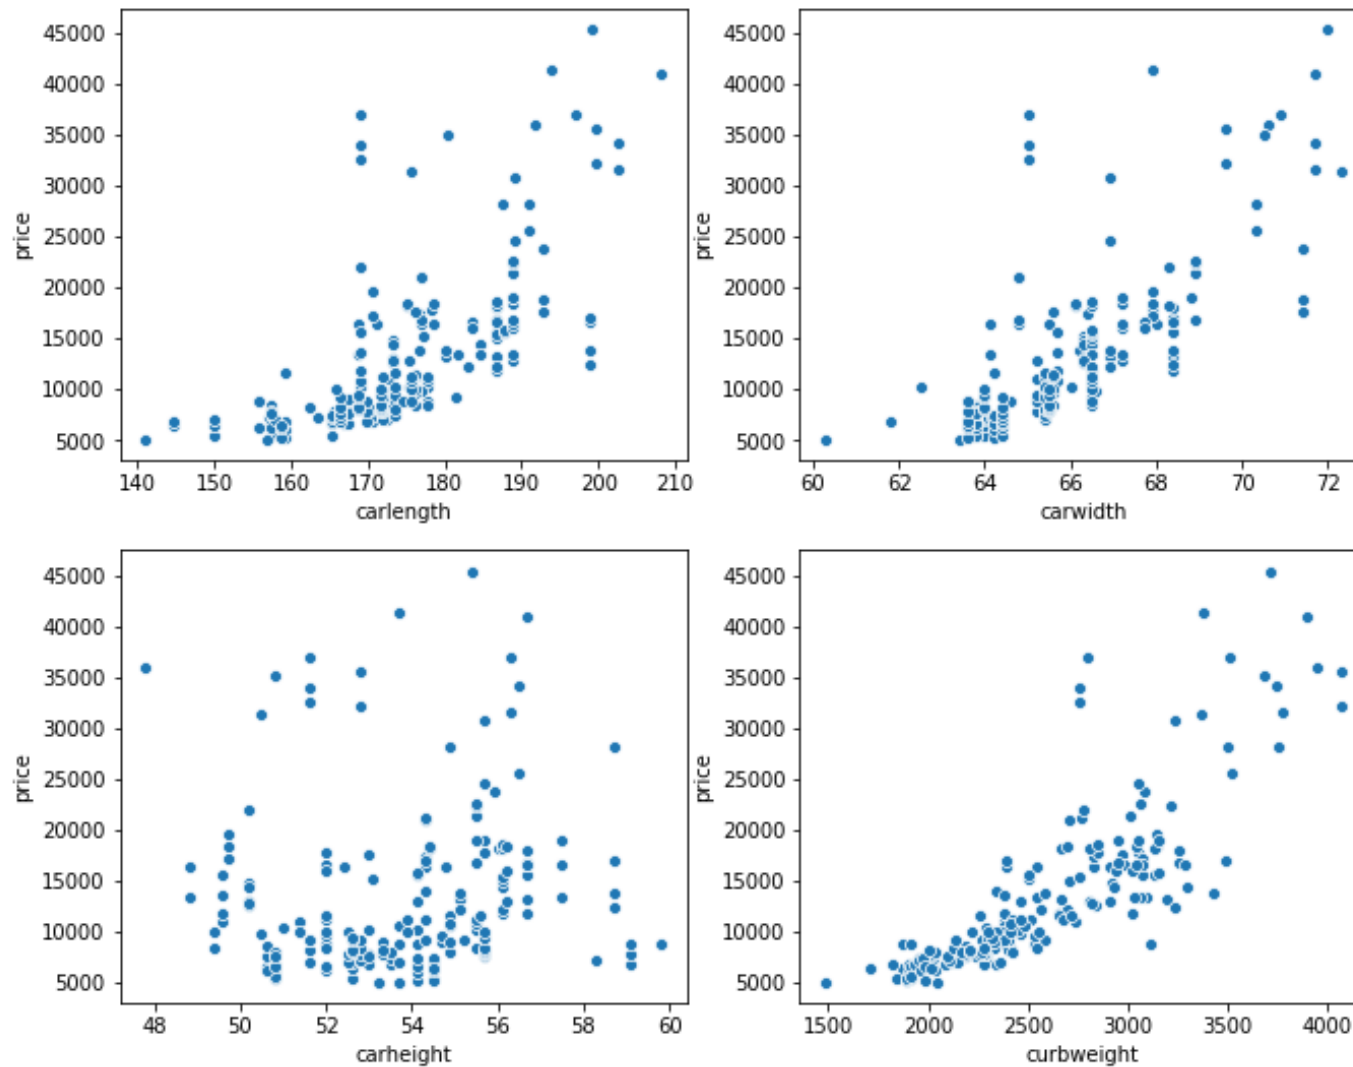
inference

- four door cars are highly preferred although they are more expensive than two door cars
- almost all the cars have front engine location and are cheaper than rear located engine cars
- most of the cars have ohc engine type but have lowest price
- dohcvt are highly expensive cars
- most of the cars have 4 cylinders



- Its assigned insurance risk rating
- A value of +3 indicates that the auto is risky, -3 that it is probably pretty safe.

Visualising numerical data

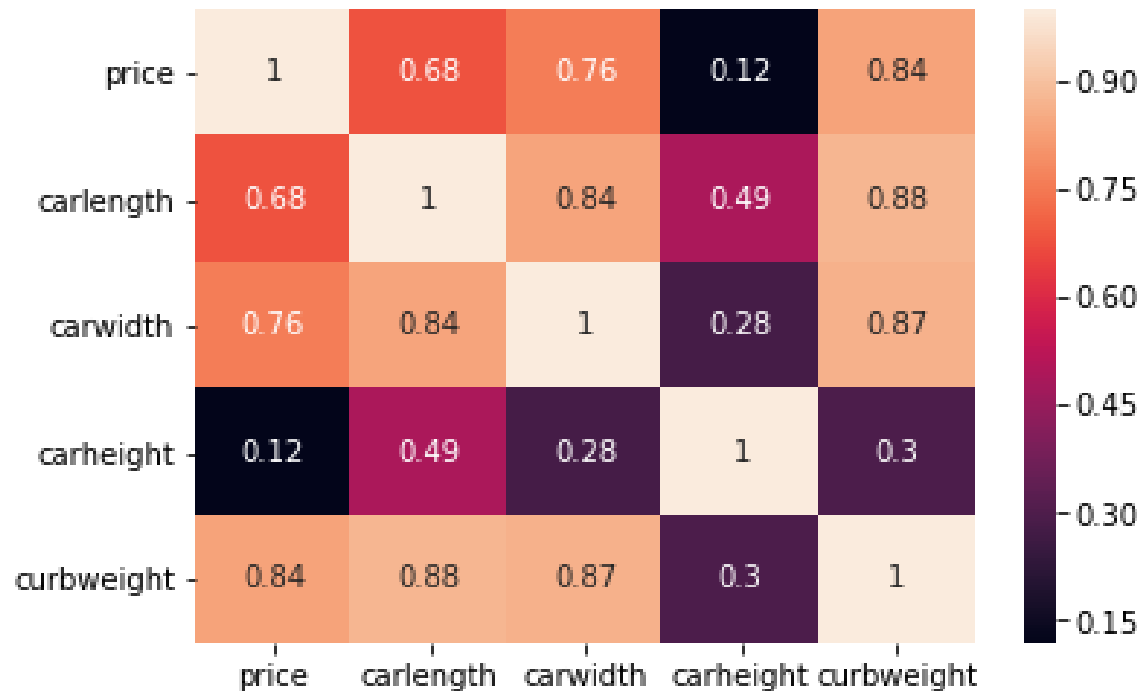


inference

- carlength, carwidth and curbweight seem to have correlation with price
- carheight does not show any significant relation with the price

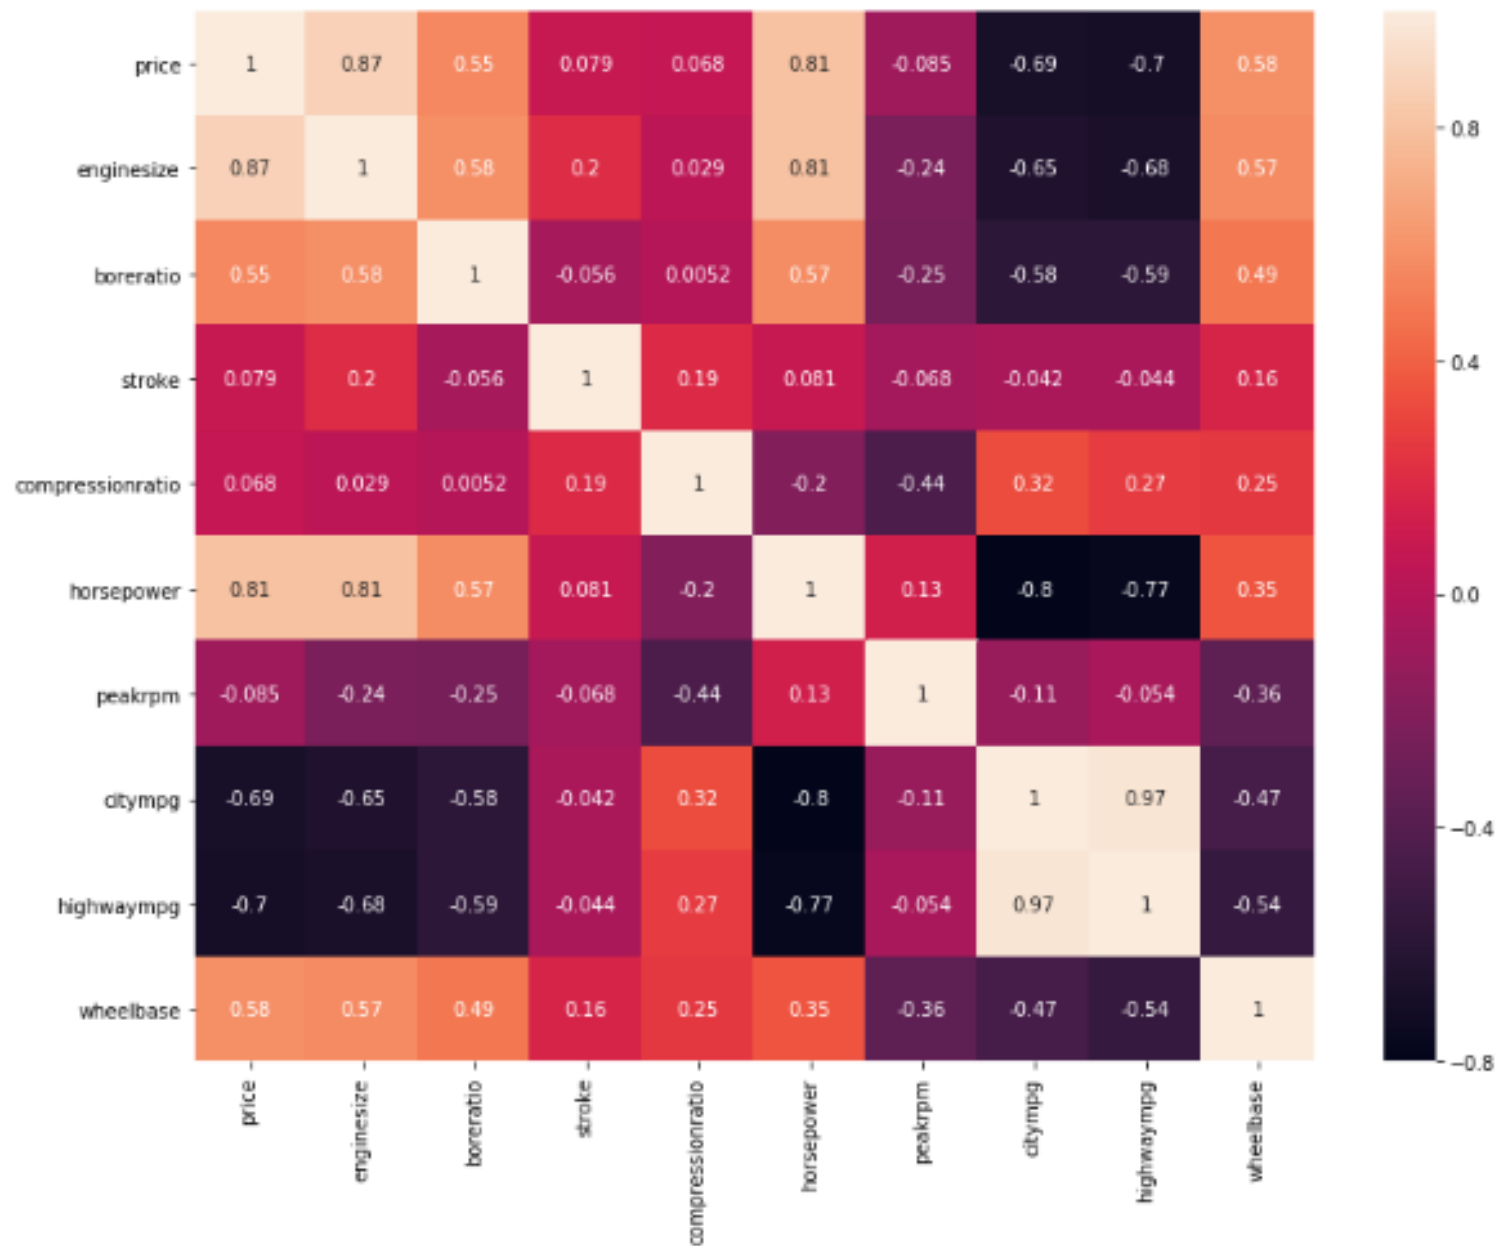
```
sns.heatmap(df_temp.corr(), annot = True)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x19033b60128>
```



inference

- price is highly correlated to curbweight



Step 4: Deriving new features

```
#Fuel economy
df['fuel_economy'] = (0.55 * df['citympg']) + (0.45 * df['highwaympg'])
```

```
# Binning the company based on average price of the car
df['price'] = df['price'].astype('int')
temp = df.copy()
table = temp.groupby(['Company'])['price'].mean()
temp = temp.merge(table.reset_index(), how='left', on='Company')
bins = [0,10000,20000,40000]
cars_bin=['Budget','Medium','Expensive']
df['carsrange'] = pd.cut(temp['price_y'],bins,right=False,labels=cars_bin)
df.head()
```

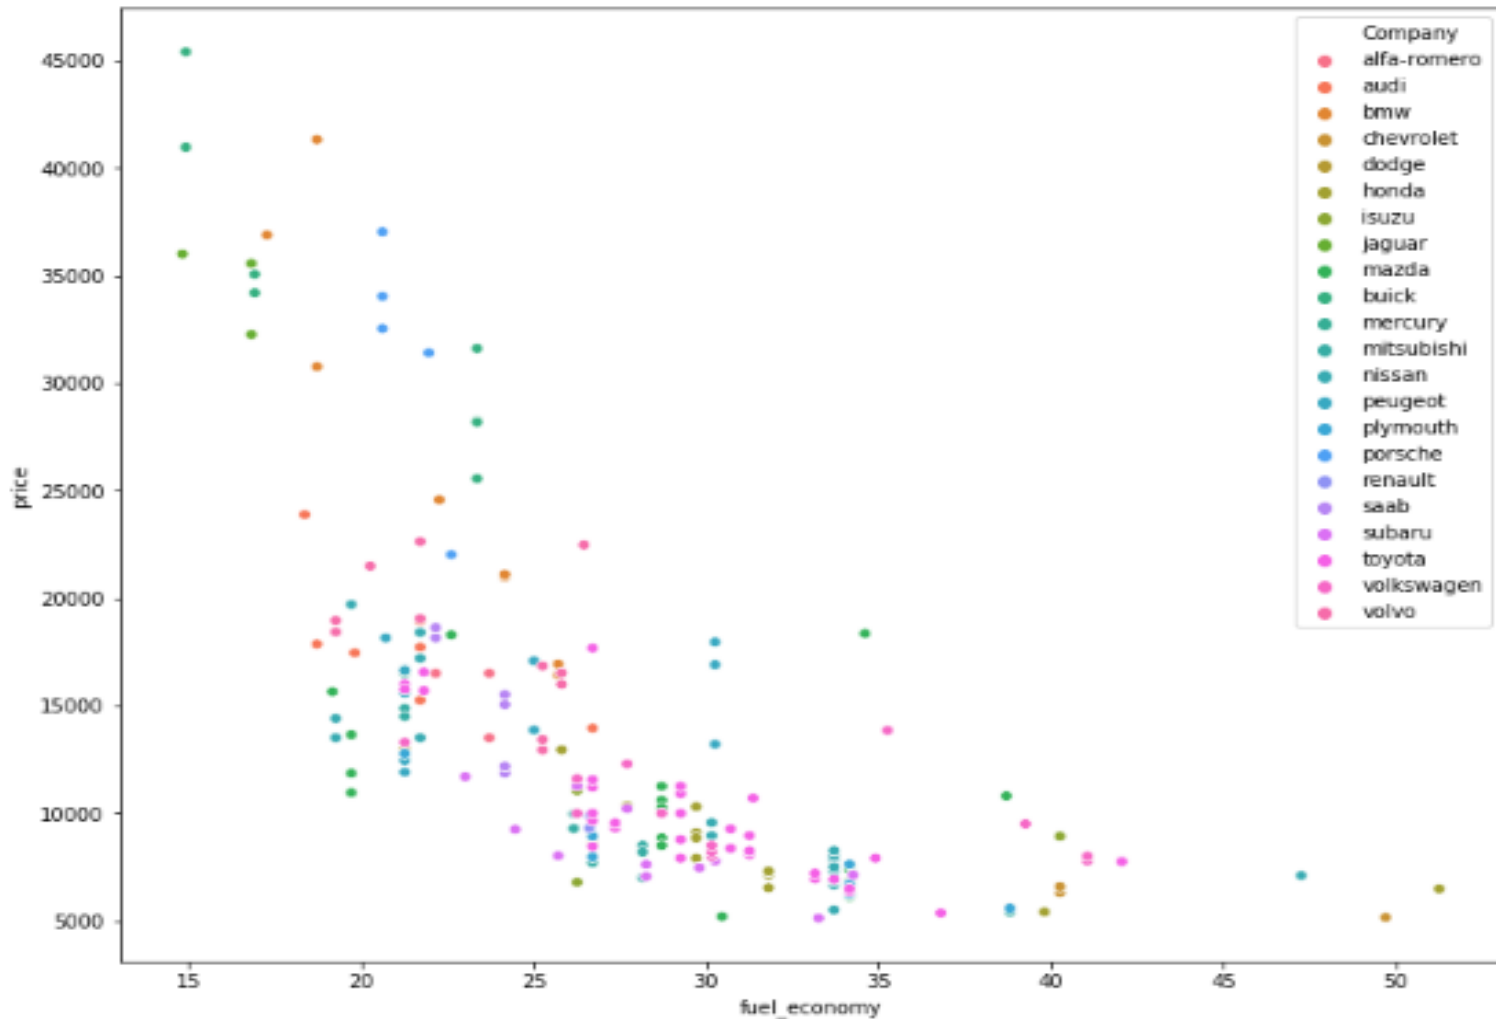
	fuel	engine	location	wheelbase	carlength	...	stroke	compressionratio	horsepower	peakrpm	citympg	highwaympg	price	Company	fuel_economy	carsrange
vd			front	88.6	168.8	...	2.68	9.0	111	5000	21	27	13495	alfa-romero	23.70	Medium
vd			front	88.6	168.8	...	2.68	9.0	111	5000	21	27	16500	alfa-romero	23.70	Medium
vd			front	94.5	171.2	...	3.47	9.0	154	5000	19	26	16500	alfa-romero	22.15	Medium
vd			front	99.8	176.6	...	3.40	10.0	102	5500	24	30	13950	audi	26.70	Medium
vd			front	99.4	176.6	...	3.40	8.0	115	5500	18	22	17450	audi	19.80	Medium

Calculating fuel_economy using citympg and highwaympg

Binning the company based on price

Dividing the car into 'Budget', 'Medium' and 'Expensive' carsrange

Step 5: Bivariate analysis



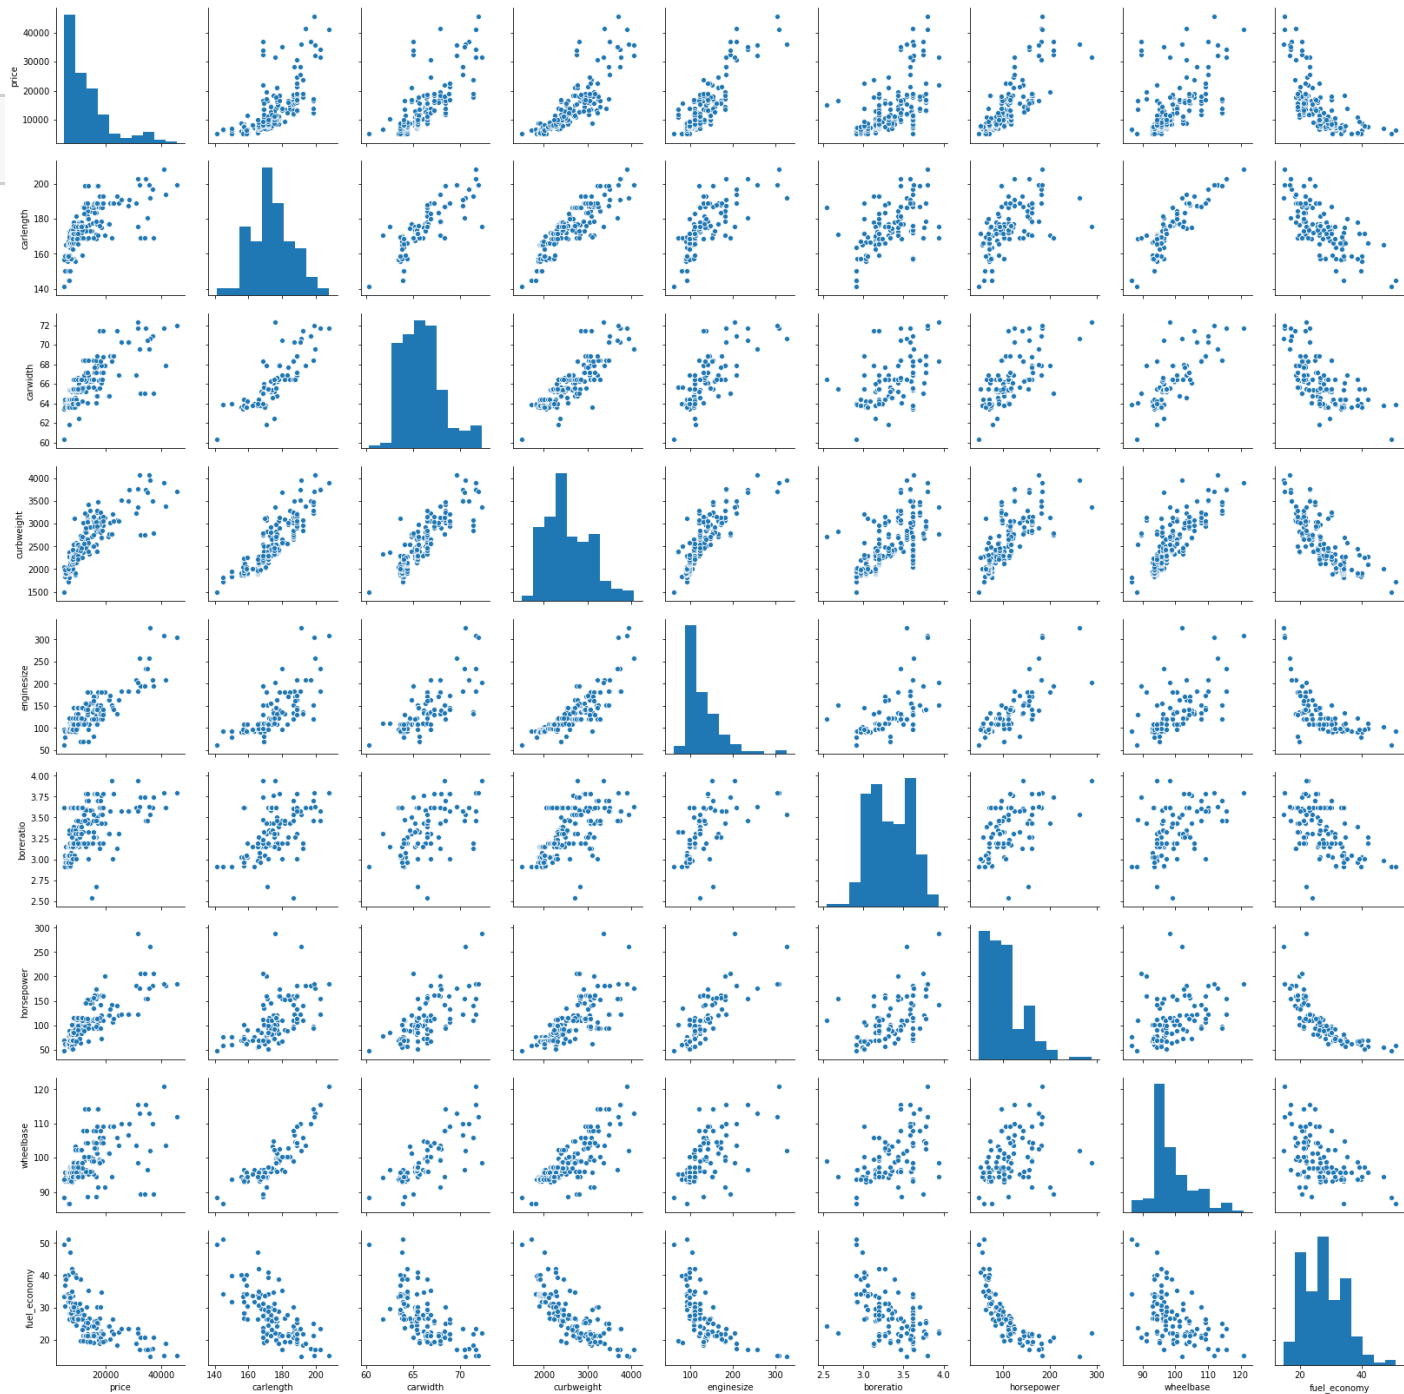
- fuel economy is negatively correlated with price

List of significant variables after analysis

- fueltype
- carbody
- drivewheel
- enginelocation
- enginetype
- cylindernumber
- aspiration
- carlength
- carwidth
- curbweight
- enginesize
- boreratio
- horsepower
- wheelbase
- fuel_economy
- carsrange

car.shape

(205, 17)



Step 6: Creating dummy variable

```
: def dummies(x,df):  
    temp = pd.get_dummies(df[x], drop_first = True)  
    df = pd.concat([df, temp], axis = 1)  
    df.drop([x], axis = 1, inplace = True)  
    return df  
  
car = dummies('carsrange',car)  
car = dummies('fueltype',car)  
car = dummies('carbody',car)  
car = dummies('drivewheel',car)  
car = dummies('engine location',car)  
car = dummies('engine type',car)  
car = dummies('cylindernumber',car)  
car = dummies('aspiration',car)
```

Creating dummy variables for the above mentioned categorical features

Step 7: Splitting the train-test data and Rescaling the features

```
: import sklearn
  from sklearn.model_selection import train_test_split

car_train, car_test = train_test_split(car, train_size = 0.8, random_state = 100)
print(car_train.shape)
print(car_test.shape)
```

(164, 32)
(41, 32)

Splitting the car DataFrame into car_train (80%) and car_test (20%)

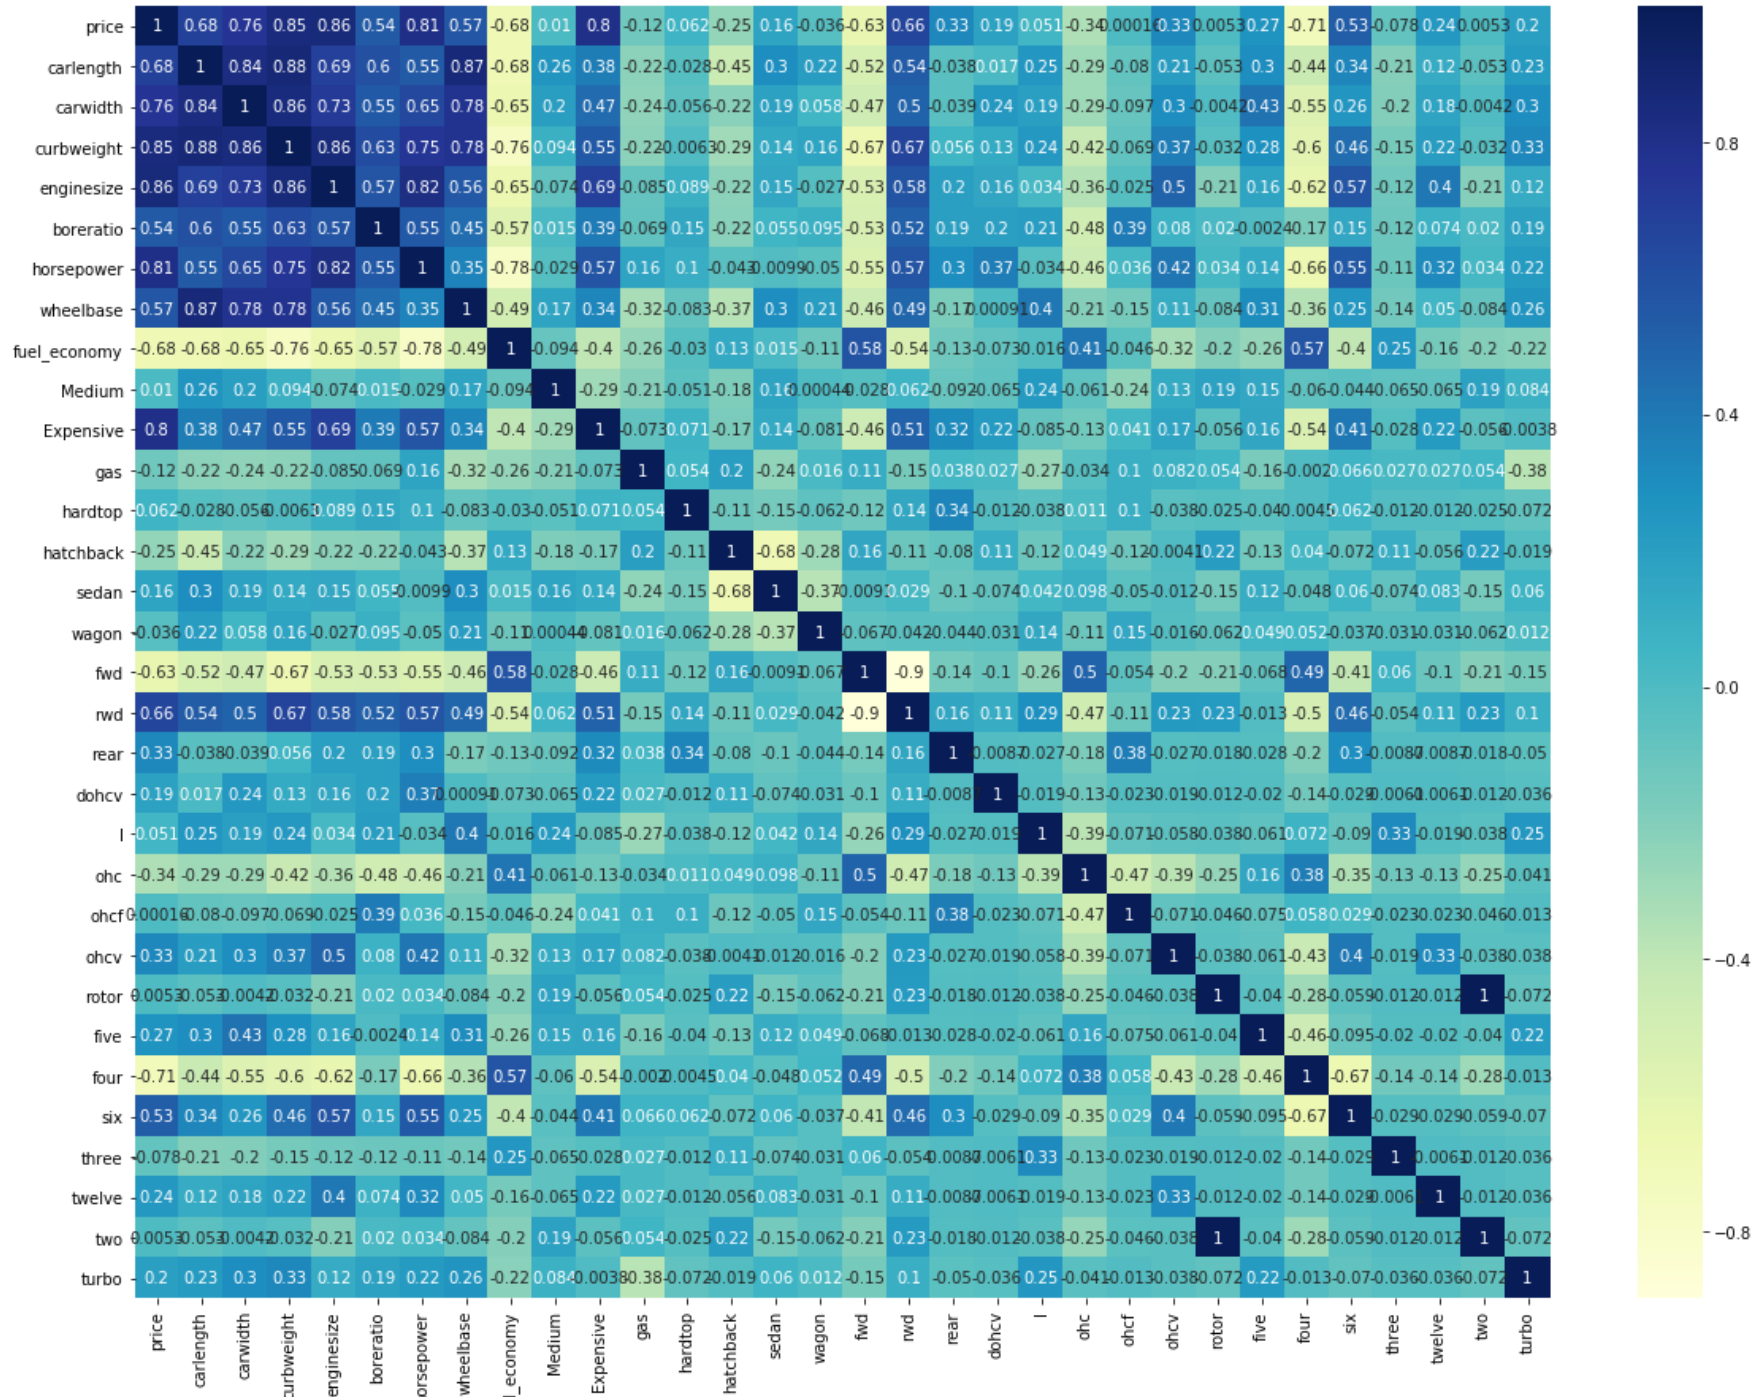
```
from sklearn.preprocessing import MinMaxScaler

# 1. instantiate an object
scaler = MinMaxScaler()

# create list of numeric vars
num_vars = ['carlength', 'carwidth', 'curbweight', 'enginesize', 'bore_ratio', 'horsepower', 'wheelbase', 'fuel_economy']

# 2. fit on data
car_train[num_vars] = scaler.fit_transform(car_train[num_vars])
car_train.head()
```

Rescaling the car_test numerical data using normalization



Step 8: Building the model

```
# dividing car_train into X and y variable  
y_train = car_train.pop('price')  
X_train = car_train
```

Using RFE(Recursive Feature Elimination) creating model and finding variables which are selected and are significant for the model

```
[('carlength', False, 12),  
 ('carwidth', True, 1),  
 ('curbweight', True, 1),  
 ('enginesize', False, 21),  
 ('boreratio', False, 11),  
 ('horsepower', True, 1),  
 ('wheelbase', False, 2),  
 ('fuel_economy', False, 3),  
 ('Medium', False, 13),  
 ('Expensive', True, 1),  
 ('gas', False, 20),  
 ('hardtop', True, 1),  
 ('hatchback', True, 1),  
 ('sedan', True, 1),  
 ('wagon', True, 1),  
 ('fwd', False, 19),  
 ('rwd', False, 15),  
 ('rear', True, 1),  
 ('dohcv', True, 1),  
 ('l', False, 17),  
 ('ohc', False, 8),  
 ('ohcf', False, 9),  
 ('ohcv', False, 10),  
 ('rotor', False, 14),  
 ('five', False, 7),  
 ('four', False, 4),  
 ('six', False, 6),  
 ('three', False, 16),  
 ('twelve', False, 5),  
 ('two', False, 18),  
 ('turbo', False, 22)]
```

Building model using statsmodel, for detailed statistics

Model 1

```
=====
                        OLS Regression Results
=====
Dep. Variable:          price      R-squared:          0.934
Model:                  OLS        Adj. R-squared:       0.930
Method:                 Least Squares    F-statistic:       216.8
Date:                  Mon, 15 Jul 2019    Prob (F-statistic): 5.55e-85
Time:                  16:10:03          Log-Likelihood:    -1476.7
No. Observations:      164             AIC:              2975.
Df Residuals:          153             BIC:              3009.
Df Model:              10
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	4350.1729	1065.491	4.083	0.000	2245.200	6455.146
carwidth	9525.1787	1877.783	5.073	0.000	5815.450	1.32e+04
curbweight	1.009e+04	2196.046	4.596	0.000	5753.520	1.44e+04
horsepower	1.077e+04	1967.349	5.475	0.000	6885.493	1.47e+04
Expensive	9133.5870	682.846	13.376	0.000	7784.563	1.05e+04
hardtop	-4119.0959	1376.708	-2.992	0.003	-6838.908	-1399.284
hatchback	-3784.5405	1029.136	-3.677	0.000	-5817.691	-1751.390
sedan	-3023.9526	1010.659	-2.992	0.003	-5020.600	-1027.305
wagon	-4304.5389	1080.327	-3.984	0.000	-6438.822	-2170.256
rear	8186.6135	1861.162	4.399	0.000	4509.720	1.19e+04
dohcv	-5948.3076	2457.501	-2.420	0.017	-1.08e+04	-1093.292

```
=====
Omnibus:              71.562    Durbin-Watson:          1.899
Prob(Omnibus):        0.000    Jarque-Bera (JB):       362.728
Skew:                 1.533    Prob(JB):               1.72e-79
Kurtosis:             9.609    Cond. No.                28.1
=====
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Model 2

OLS Regression Results

Dep. Variable:	price	R-squared:	0.930			
Model:	OLS	Adj. R-squared:	0.926			
Method:	Least Squares	F-statistic:	228.1			
Date:	Mon, 15 Jul 2019	Prob (F-statistic):	2.72e-84			
Time:	16:10:04	Log-Likelihood:	-1481.3			
No. Observations:	164	AIC:	2983.			
Df Residuals:	154	BIC:	3014.			
Df Model:	9					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	2557.2278	903.465	2.830	0.005	772.443	4342.013
carwidth	9248.6200	1923.319	4.809	0.000	5449.126	1.3e+04
curbweight	1.049e+04	2247.943	4.665	0.000	6046.897	1.49e+04
horsepower	1.061e+04	2016.756	5.262	0.000	6627.942	1.46e+04
Expensive	9236.6506	699.362	13.207	0.000	7855.068	1.06e+04
hatchback	-1976.6593	854.312	-2.314	0.022	-3664.343	-288.975
sedan	-1245.5402	838.230	-1.486	0.139	-2901.454	410.373
wagon	-2542.0795	928.689	-2.737	0.007	-4376.693	-707.466
rear	7833.3739	1904.765	4.113	0.000	4070.534	1.16e+04
dohcv	-5917.8372	2520.129	-2.348	0.020	-1.09e+04	-939.352
=====						
Omnibus:	60.823	Durbin-Watson:	1.818			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	260.864			
Skew:	1.324	Prob(JB):	2.26e-57			
Kurtosis:	8.582	Cond. No.	28.0			
=====						

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Calculating VIF

	Features	VIF
0	const	30.64
2	curbweight	7.72
6	sedan	6.57
5	hatchback	6.16
1	carwidth	4.47
3	horsepower	4.08
7	wagon	3.76
4	Expensive	1.79
8	rear	1.64
9	dohcv	1.44

VIF > 5 should not be ignored
therefore dropping feature "curbweight" as it is
showing high multicollinearity

Model 3

OLS Regression Results

```

=====
Dep. Variable:          price    R-squared:          0.920
Model:                  OLS      Adj. R-squared:       0.916
Method:                 Least Squares    F-statistic:       223.9
Date:                   Mon, 15 Jul 2019    Prob (F-statistic): 4.46e-81
Time:                   16:10:04    Log-Likelihood:    -1492.2
No. Observations:       164    AIC:               3002.
Df Residuals:           155    BIC:               3030.
Df Model:                8
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	2880.4237	959.253	3.003	0.003	985.527	4775.320
carwidth	1.557e+04	1454.293	10.704	0.000	1.27e+04	1.84e+04
horsepower	1.613e+04	1738.977	9.277	0.000	1.27e+04	1.96e+04
Expensive	1.009e+04	718.903	14.033	0.000	8668.525	1.15e+04
hatchback	-2617.7641	897.898	-2.915	0.004	-4391.460	-844.069
sedan	-1525.3648	890.330	-1.713	0.089	-3284.111	233.381
wagon	-1988.7744	980.849	-2.028	0.044	-3926.330	-51.219
rear	5773.5472	1973.113	2.926	0.004	1875.886	9671.209
dohcv	-1.065e+04	2456.529	-4.336	0.000	-1.55e+04	-5798.987

```

=====
Omnibus:                50.533    Durbin-Watson:       1.967
Prob(Omnibus):           0.000    Jarque-Bera (JB):    163.455
Skew:                    1.177    Prob(JB):            3.21e-36
Kurtosis:                7.287    Cond. No.            19.5
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Model 4

OLS Regression Results

```

=====
Dep. Variable:          price    R-squared:          0.916
Model:                  OLS      Adj. R-squared:       0.912
Method:                 Least Squares    F-statistic:        242.9
Date:                   Mon, 15 Jul 2019    Prob (F-statistic):  1.75e-80
Time:                   16:10:04    Log-Likelihood:     -1496.6
No. Observations:       164    AIC:                3009.
Df Residuals:           156    BIC:                3034.
Df Model:                7
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	4026.4943	896.622	4.491	0.000	2255.409	5797.580
carwidth	1.429e+04	1420.152	10.060	0.000	1.15e+04	1.71e+04
horsepower	1.762e+04	1703.238	10.343	0.000	1.43e+04	2.1e+04
Expensive	1.056e+04	717.319	14.723	0.000	9144.129	1.2e+04
hatchback	-3565.6485	857.484	-4.158	0.000	-5259.425	-1871.872
sedan	-2442.3525	853.314	-2.862	0.005	-4127.892	-756.813
wagon	-2837.5598	959.416	-2.958	0.004	-4732.682	-942.437
dohcv	-1.153e+04	2496.688	-4.616	0.000	-1.65e+04	-6594.222

```

=====
Omnibus:                 38.175    Durbin-Watson:         2.023
Prob(Omnibus):           0.000    Jarque-Bera (JB):      101.397
Skew:                    0.938    Prob(JB):              9.59e-23
Kurtosis:                6.365    Cond. No.              18.7
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
#Calculating the Variance Inflation Factor  
checkVIF(X_train_new)
```

	Features	VIF
0	const	25.38
5	sedan	5.73
4	hatchback	5.22
6	wagon	3.38
2	horsepower	2.45
1	carwidth	2.05
3	Expensive	1.59
7	dohcv	1.19

VIF > 5 should not be ignored
therefore dropping feature "sedan" as it is showing high
multicollinearity

Model 5

OLS Regression Results

Dep. Variable:	price	R-squared:	0.912
Model:	OLS	Adj. R-squared:	0.908
Method:	Least Squares	F-statistic:	269.7
Date:	Mon, 15 Jul 2019	Prob (F-statistic):	5.48e-80
Time:	16:10:05	Log-Likelihood:	-1500.8
No. Observations:	164	AIC:	3016.
Df Residuals:	157	BIC:	3037.
Df Model:	6		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	2063.0393	590.433	3.494	0.001	896.823	3229.256
carwidth	1.313e+04	1392.418	9.431	0.000	1.04e+04	1.59e+04
horsepower	1.894e+04	1676.599	11.296	0.000	1.56e+04	2.23e+04
Expensive	1.072e+04	731.419	14.654	0.000	9273.261	1.22e+04
hatchback	-1414.9357	422.449	-3.349	0.001	-2249.352	-580.519
wagon	-590.0016	563.712	-1.047	0.297	-1703.440	523.436
dohcv	-1.204e+04	2546.715	-4.726	0.000	-1.71e+04	-7005.974

Omnibus:	32.476	Durbin-Watson:	2.034
Prob(Omnibus):	0.000	Jarque-Bera (JB):	66.183
Skew:	0.896	Prob(JB):	4.25e-15
Kurtosis:	5.545	Cond. No.	17.1

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Model 6

OLS Regression Results

Dep. Variable:	price	R-squared:	0.911			
Model:	OLS	Adj. R-squared:	0.908			
Method:	Least Squares	F-statistic:	323.2			
Date:	Mon, 15 Jul 2019	Prob (F-statistic):	4.98e-81			
Time:	16:10:05	Log-Likelihood:	-1501.3			
No. Observations:	164	AIC:	3015.			
Df Residuals:	158	BIC:	3033.			
Df Model:	5					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	1962.6110	582.760	3.368	0.001	811.607	3113.615
carwidth	1.304e+04	1389.831	9.379	0.000	1.03e+04	1.58e+04
horsepower	1.899e+04	1676.369	11.328	0.000	1.57e+04	2.23e+04
Expensive	1.082e+04	725.477	14.910	0.000	9384.228	1.22e+04
hatchback	-1289.2380	405.141	-3.182	0.002	-2089.428	-489.048
dohcv	-1.212e+04	2546.316	-4.759	0.000	-1.71e+04	-7087.720
=====						
Omnibus:	34.628	Durbin-Watson:	2.031			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	71.778			
Skew:	0.947	Prob(JB):	2.59e-16			
Kurtosis:	5.630	Cond. No.	17.0			
=====						

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

	Features	VIF
0	const	10.25
2	horsepower	2.27
1	carwidth	1.88
3	Expensive	1.55
5	dohcv	1.19
4	hatchback	1.11

Model looks fine as all the values are less than 5

dropping "hatchback" just to check the statistics of the model

Model 7

	Features	VIF
0	const	8.49
2	horsepower	2.22
1	carwidth	1.78
3	Expensive	1.52
4	dohcv	1.16

```

OLS Regression Results

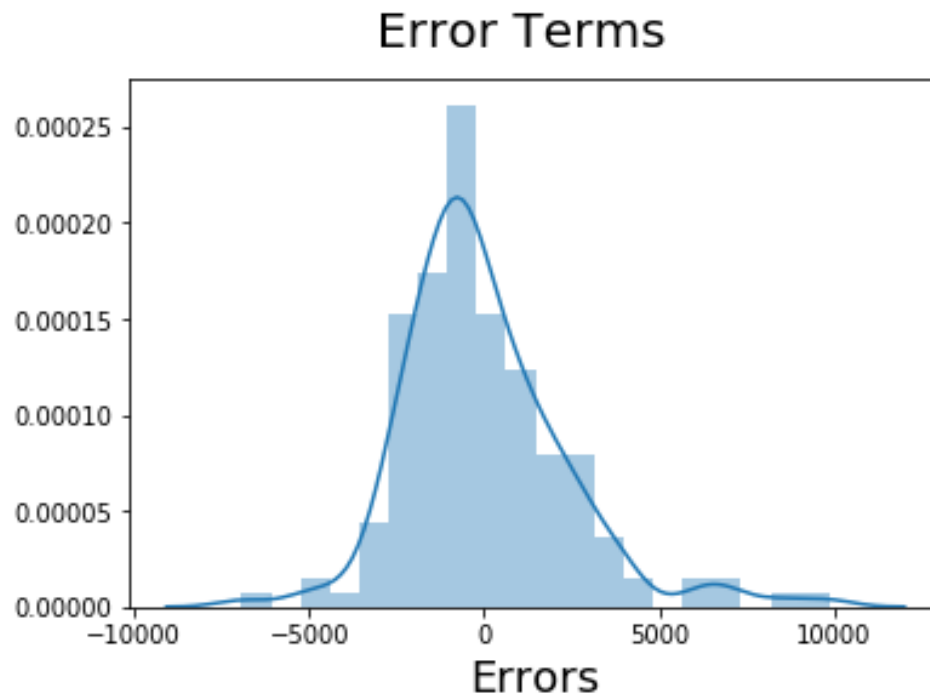
=====
Dep. Variable:      price      R-squared:      0.905
Model:              OLS       Adj. R-squared:  0.903
Method:             Least Squares   F-statistic:    379.7
Date:               Mon, 15 Jul 2019   Prob (F-statistic): 3.22e-80
Time:               16:10:05    Log-Likelihood: -1506.4
No. Observations:   164          AIC:             3023.
Df Residuals:       159          BIC:             3038.
Df Model:           4
Covariance Type:    nonrobust
=====
                    coef      std err          t      P>|t|      [0.025      0.975]
-----
const             1193.7288      545.317        2.189      0.030      116.730      2270.728
carwidth          1.405e+04      1391.088       10.100      0.000      1.13e+04      1.68e+04
horsepower        1.825e+04      1707.143       10.691      0.000      1.49e+04      2.16e+04
Expensive         1.116e+04       737.543       15.137      0.000      9707.228      1.26e+04
dohcv             -1.326e+04      2592.276       -5.114      0.000     -1.84e+04     -8138.382
=====
Omnibus:           39.069    Durbin-Watson:      2.010
Prob(Omnibus):     0.000    Jarque-Bera (JB):   84.029
Skew:              1.052    Prob(JB):           5.67e-19
Kurtosis:          5.806    Cond. No.           16.1
=====

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```

Final model as the p-value is less than 0.05 and VIF is also less than 5

Step 9: Residual analysis on train set



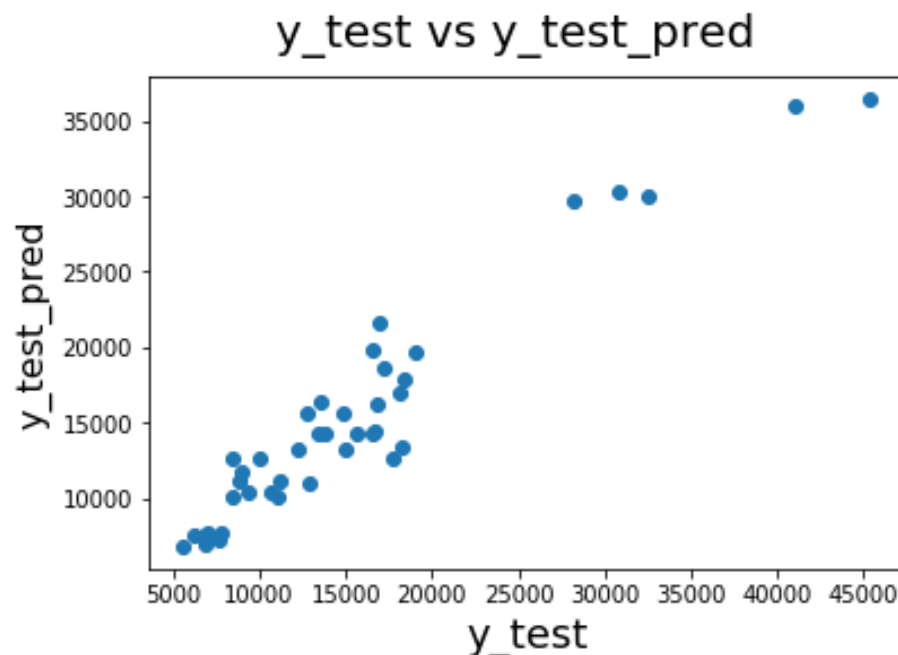
- errors seem to be normally distributed with mean = 0

Step 10: Prediction and evaluation

```
# making prediction  
y_test_pred = lm.predict(X_test_new)
```

```
# evaluation  
from sklearn.metrics import r2_score  
r2_score(y_test, y_test_pred)
```

0.9094330812753356



OLS Regression Results

```

=====
Dep. Variable:          price    R-squared:          0.905
Model:                  OLS      Adj. R-squared:       0.903
Method:                 Least Squares    F-statistic:        379.7
Date:                  Mon, 15 Jul 2019    Prob (F-statistic):  3.22e-80
Time:                  16:10:07    Log-Likelihood:     -1506.4
No. Observations:      164    AIC:                3023.
Df Residuals:          159    BIC:                3038.
Df Model:               4
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	1193.7288	545.317	2.189	0.030	116.730	2270.728
carwidth	1.405e+04	1391.088	10.100	0.000	1.13e+04	1.68e+04
horsepower	1.825e+04	1707.143	10.691	0.000	1.49e+04	2.16e+04
Expensive	1.116e+04	737.543	15.137	0.000	9707.228	1.26e+04
dohcv	-1.326e+04	2592.276	-5.114	0.000	-1.84e+04	-8138.382

```

=====
Omnibus:                39.069    Durbin-Watson:        2.010
Prob(Omnibus):          0.000    Jarque-Bera (JB):     84.029
Skew:                   1.052    Prob(JB):             5.67e-19
Kurtosis:               5.806    Cond. No.             16.1
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

since R-squared and adjusted R-squared are 0.912 and 0.909, hence 90% variance explained by the model

all the p-values are less than 0.05, therefore we can say the coefficients are statistically significant

